

# Personal Review Overview

Katharina Boudgoust

December 21, 2025

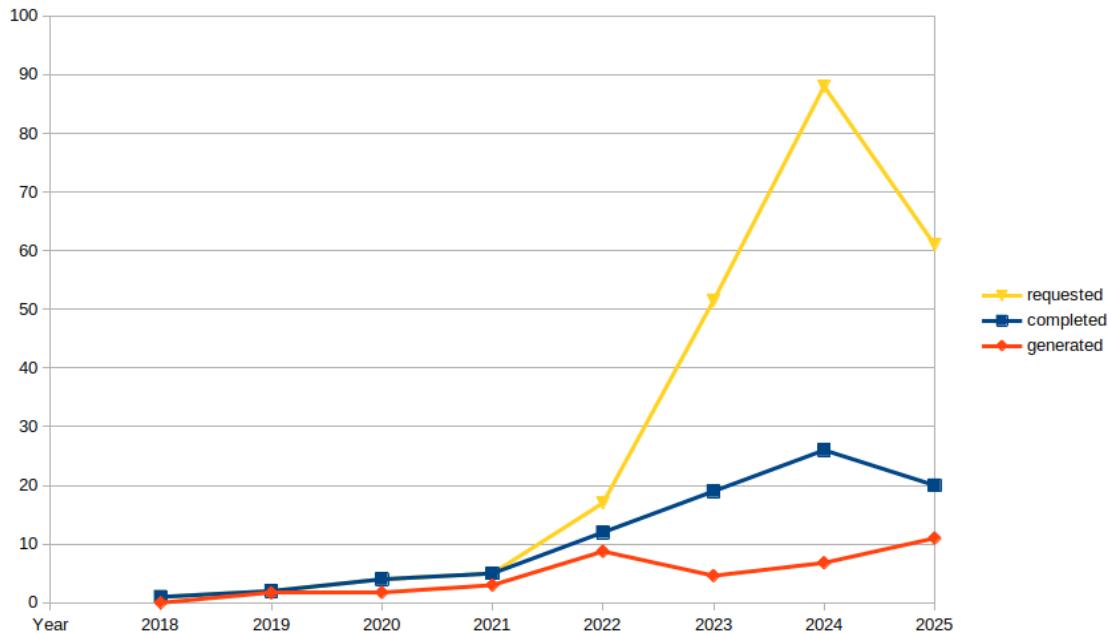


Figure 1: My history on requested, completed and generated reviews.

**Context.** Inspired by the review documentation by my colleague [Alice Pellet–Mary](#), I decided in 2023 to follow her and also document my personal review history. By making it publicly available on my website, I share it in the hope that it might be insightful or inspiring for others.

This personal review history has multiple roles. First, it simply documents my personal review load. Given the numerous tasks a researcher is occupied with during the year, I find it very valuable to keep track of some of my tasks. Second, it helps me to take a step aback and decide, whenever I receive a new review request or program committee (PC) invitation, whether I want to accept or decline it. For instance, one can fix an a-priori ratio  $\tau$  between *completed* and *generated* reviews one wants to achieve throughout a year, for instance 3:1 or 4:1, and start declining once the ratio is reached.

My review history in Figure 1 shows that during my PhD (2018-2021) completing and generating reviews was pretty much in balance ( $\tau < 2.3$ ). Starting from my Postdoc (2022-2023), I was invited to program committees for conferences, but also more and more requested as subreviewer for other PC members and as reviewer for journals. On the other hand, my own productivity went down, as I couldn't spend as much time as during my PhD on "simply" doing research (in 2023,  $\tau = 4.75$ ). I quickly learnt to decline reviews or distribute them to subreviewers (in 2024,  $\tau = 3.82$  and in 2025,  $\tau = 1.82$ ).

The number of requested reviews went significantly down in 2025. Without having evidence for it, I wish to believe that communicating about my review (over)load helped raising consciousness and let people think twice before asking me for a review.

**Requested review.** An article I was asked/assigned to review. It covers the following three cases:

- Case 1) I was requested by an editor of a journal,
- Case 2) A program committee member of a conference asked me to do a subreview,
- Case 3) The papers that got assigned to me as member of a conference program committee or
- Case 4) The estimated paper load for a declined program committee invitation.

The review counts for the year in which I actually wrote it. For declined program committee invitations, I do the following (lower-bound) estimations of the review load.<sup>1</sup> I count 15 papers for big PC duties (e.g. Eurocrypt, Crypto, Asiacrypt, TCC), 10 papers for medium PC duties (e.g. PKC, CCS) and 5 papers for small PC duties (e.g. Latincrypt, PQCrypto, FHE.org).

**Completed review.** An article I personally read and reviewed. Journal reviews are counted once, even when there are multiple iterations. Articles that I reviewed multiple times for different venues are counted multiple times. The difference between *requested* and *completed* reviews thus contains articles I declined to review (Case 1 and Case 2), I delegated to subreviewers to review it (Case 3) or PC invitations I declined (Case 4).

**Generated review.** For each submission, this is the number of reviews obtained for the submission divided by the number of co-authors on the article.

---

<sup>1</sup>This estimation is novel since 2025. In earlier versions of this document, there was no Case 4. Thus, I did not count the reviews that would have come with the duties of being a program committee member, if I declined the invitation for the program committee.